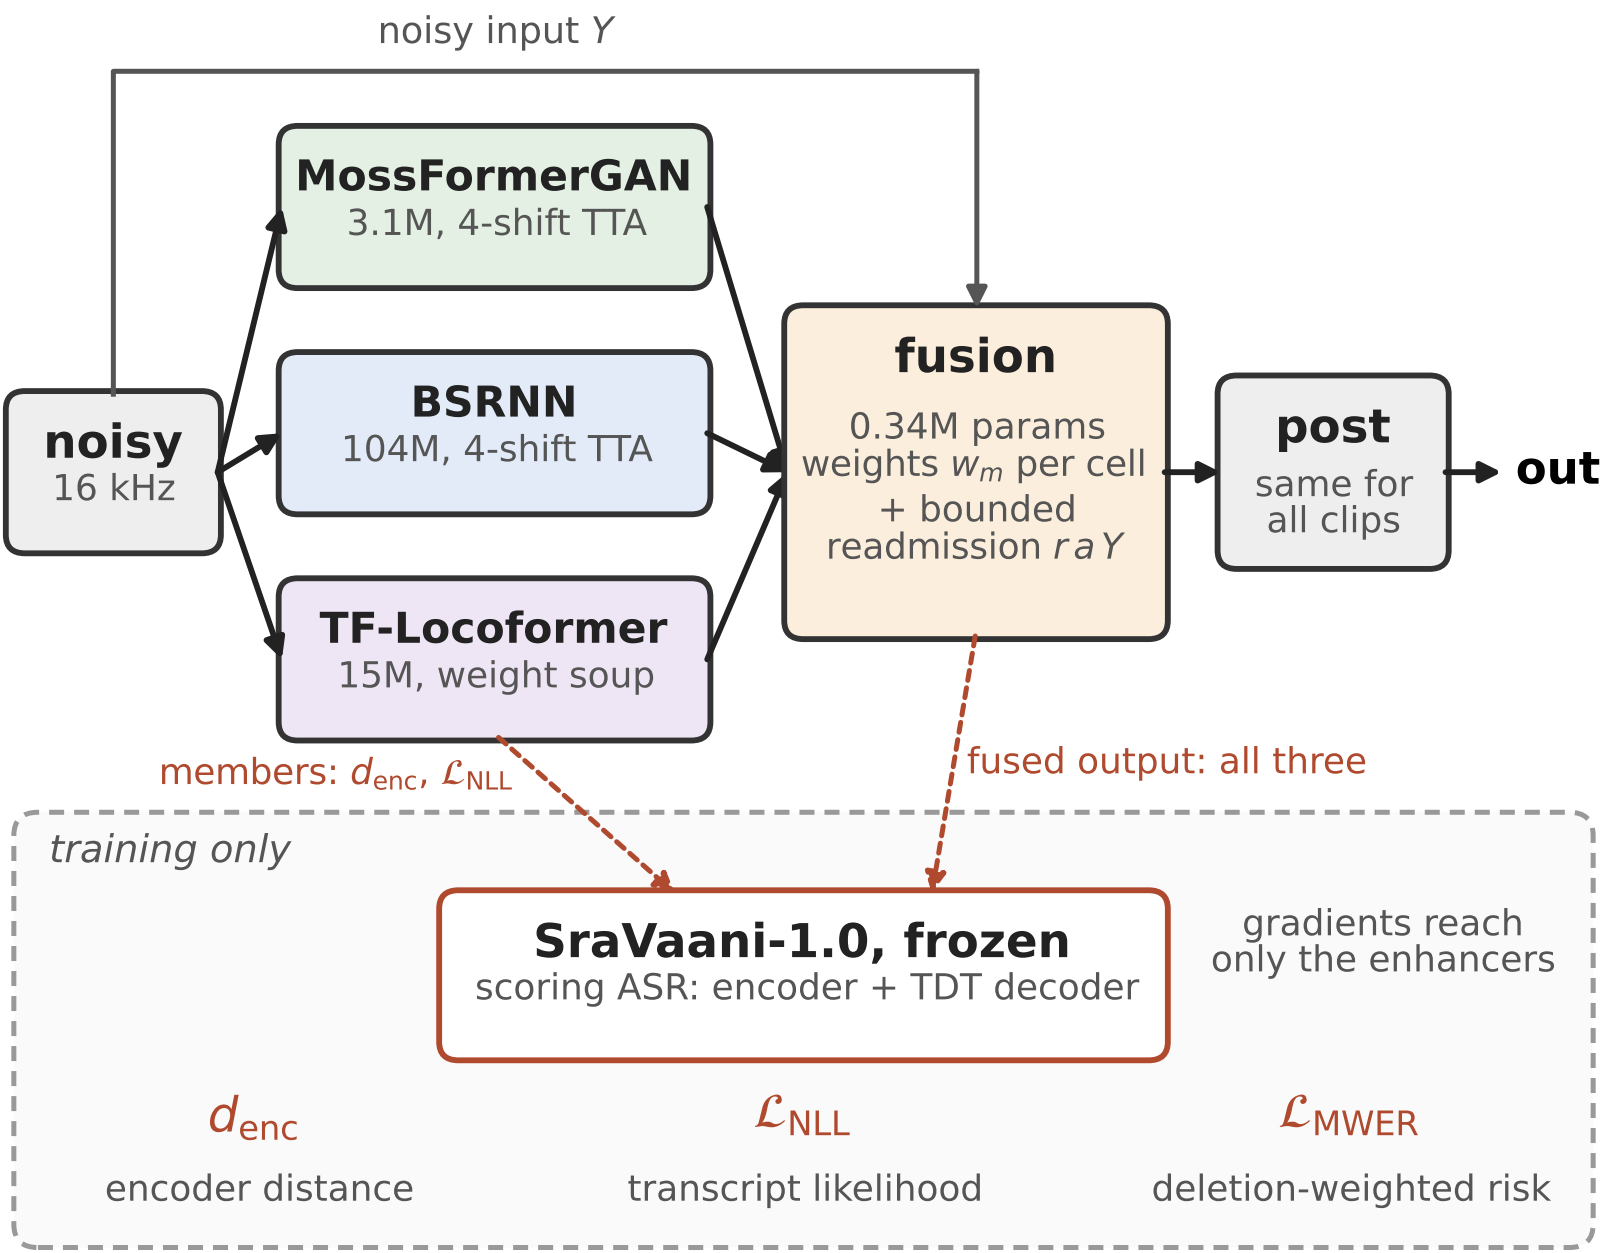
noisy input Y
MossFormerGAN
3.1M, 4-shift TTA
BSRNN
104M, 4-shift TTA
TF-Locoformer
15M, weight soup
noisy
16 kHz
fusion
0.34M params
weights w_m per cell
+ bounded
readmission $r\,a\,Y$
post
same for
all clips
out
members: d_{enc}, $\mathcal{L}_{\text{NLL}}$
fused output: all three
training only
SraVaani-1.0, frozen
scoring ASR: encoder + TDT decoder
gradients reach
only the enhancers
d_{enc}
encoder distance
$\mathcal{L}_{\text{NLL}}$
transcript likelihood
$\mathcal{L}_{\text{MWER}}$
deletion-weighted risk